

Computational Imaging



Lecture 12: Computing Toolbox:
Image Gradient and Processing II



Qilin Sun (孙启霖)

School of Data Science

The Chinese University of Hong Kong, Shenzhen



Today's Topic

- Efficient Poisson Solver
- Poisson Image Editing
- Flash/no-flash Photography
- Gradient Cameras
 - Diff Sensor
 - Diff Optics
 - Diff Time

Efficient Poisson Solver



The Optimization Problem

Variational problem

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \quad \text{with} \quad f|_{\partial\Omega} = f^*|_{\partial\Omega}$$

gradient of f looks
like vector field \mathbf{v}

f is equivalent to
 f^* at the
boundaries

Input vector field:
 $\mathbf{v} = (u, v)$

Recall ...

Nabla operator definition

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

And for discrete images:

partial-x derivative filter

1	-1
---	----

partial-y derivative filter

1
-1



The Optimization Problem

We can use the gradient approximation to discretize the variational problem

Discrete problem

What are G , f , and v ?

$$\min_f \|Gf - v\|^2 \xrightarrow{\text{图像}} \min_f \|\nabla f - v\|^2$$

We will ignore the boundary conditions for now.

Recall ...

Nabla operator definition

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

And for discrete images:

partial-x derivative filter

1	-1
---	----

partial-y derivative filter

1
-1



The Optimization Problem

We can use the gradient approximation to discretize the variational problem

Discrete problem

matrix G formed by stacking together discrete gradients

$$\min_f \|Gf - v\|^2$$

vectorized version of the unknown image

vectorized version of the target gradient field

We will ignore the boundary conditions for now.

How to solve this optimization problem?

Recall ...

Nabla operator definition

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

And for discrete images:

partial-x derivative filter

1	-1
---	----

partial-y derivative filter

1
-1



Solution 1: Compute Stationary Points

Given the loss function:

$$E(f) = \|Gf - v\|^2$$

Compute its derivative:

$$\frac{\partial E}{\partial f} = G^T G f - G^T v$$

Set that to zero:

$$\frac{\partial E}{\partial f} = 0 \Rightarrow \underbrace{G^T G f}_{\text{—}} = \underbrace{G^T v}_{\text{—}}$$

It is equal to the vector b we derived previously!

It is equal to the Laplacian matrix A we derived previously!



Solution 2: Gradient Descent

Given the loss function:

$$E(f) = \|Gf - v\|^2$$

Compute its derivative:

$$\frac{\partial E}{\partial f} = G^T G f - G^T v = \boxed{Af - b \equiv -r}$$

The *residual*

Iteratively compute a solution:

$$f^{i+1} = f^i + \eta^i r^i \quad \text{for } i = 0, 1, \dots, N, \\ \text{where } \eta^i \text{ are positive step sizes}$$



Solution 2: Optimal Step Sizes

Make derivative of loss function *with respect to η^i* equal to zero:

$$E(f) = \|Gf - v\|^2$$

$$E(f^{i+1}) = \|G(f^i + \eta^i r^i) - v\|^2$$

$$\frac{\partial E(f^{i+1})}{\partial \eta^i} = [b - A(f^i + \eta^i r^i)]^T r^i = 0 \Rightarrow \eta^i = \frac{(r^i)^T r^i}{(r^i)^T A r^i}$$



Solution 2: Gradient Descent

Given the loss function:

$$E(f) = \|Gf - v\|^2$$

Minimize by iteratively computing:

$$r^i = b - \boxed{Af^i}, \quad \eta^i = \frac{(r^i)^T r^i}{(r^i)^T Ar^i}, \quad f^{i+1} = f^i + \eta^i r^i, \quad i = 0, \dots, N$$

Is this cheaper than the pseudo-inverse approach?

- We never need to compute \mathbf{A} , only its products with vectors \mathbf{f}, \mathbf{r} .
- Vectors \mathbf{f}, \mathbf{r} are images.
- Because \mathbf{A} is the *Laplacian matrix*, these matrix-vector products can be efficiently computed using *convolutions* with the *Laplacian kernel*.



Poractical Solution: Conjugate Gradient (CG)

Given the loss function:

Minimize by iteratively computing:

$$d^i = r^i + \beta^i d^i, \quad \eta^i = \frac{(r^i)^T r^i}{(d^i)^T Ad^i}, \quad f^{i+1} = f^i + \eta^i d^i, \quad i = 0, \dots, N$$
$$r^{i+1} = r^i - \eta [Ad^i], \quad \beta^i = \frac{(r^{i+1})^T r^{i+1}}{(r^i)^T r^i}$$

- Smarter way for selecting **update** directions
- Everything can still be done using **convolutions**
- Only **one** convolution needed per iteration



Poractical Solution: CG Initialization

Does the initialization f^0 matter?

- It doesn't matter. The loss is convex.

$$E(f) = \|Gf - v\|^2$$

It does matter in terms of **convergence speed**.

- We can use a *multi-scale* approach (like mipmap):
 - Solve an initial problem for a very low-resolution f (e.g., 2×2).
 - Use the solution to initialize gradient descent for a higher resolution f (e.g., 4×4).
 - ...
 - Use the solution to initialize gradient descent for an f with the original resolution $P \times P$.
- *Multi-grid* algorithms alternative between higher and lower resolutions during the (conjugate) gradient descent iterative procedure.



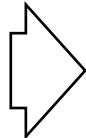
Reminder from Variational Case

Poisson equation (with Dirichlet boundary conditions)

$$\Delta f = \operatorname{div} \mathbf{v} \quad \text{over } \Omega, \quad \text{with } f|_{\partial\Omega} = f^*|_{\partial\Omega}$$

After discretization, equivalent to:

$$\begin{bmatrix} D & I & 0 & 0 & 0 & \cdots & 0 \\ I & D & I & 0 & 0 & \cdots & 0 \\ 0 & I & D & I & 0 & \cdots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & \cdots & 0 & I & D & I & 0 \\ 0 & \cdots & \cdots & 0 & I & D & I \\ 0 & \cdots & \cdots & \cdots & 0 & I & D \end{bmatrix} \cdot \begin{bmatrix} f_1 \\ f_{q_1} \\ \vdots \\ f_{q_2} \\ f_p \\ f_{q_3} \\ \vdots \\ f_{q_4} \\ \vdots \\ f_P \end{bmatrix} = \begin{bmatrix} (\nabla \cdot \mathbf{v})_1 \\ \vdots \\ (\nabla \cdot \mathbf{v})_{q_1} \\ \vdots \\ (\nabla \cdot \mathbf{v})_{q_2} \\ (\nabla \cdot \mathbf{v})_p \\ (\nabla \cdot \mathbf{v})_{q_3} \\ \vdots \\ (\nabla \cdot \mathbf{v})_{q_4} \\ \vdots \\ (\nabla \cdot \mathbf{v})_P \end{bmatrix}$$



Linear system of equations:

$$Af = b$$

Remember that what we are doing is equivalent to solving this linear system.



Note: Preconditioning

We are solving this linear system:

$$Af = b$$

For any **invertible matrix P**, this is equivalent to solving:

$$P^{-1}Af = P^{-1}b$$

Preconditioning can be incorporated in the conjugate gradient descent algorithm.

When is it preferable to solve this alternative linear system?

- Ideally: If A is invertible, and P is the same as A, the linear system becomes trivial! But computing the inverse of A is even more expensive than solving the original linear system.
- In practice: If the matrix $P^{-1}A$ has a better condition number, or its singular values are more uniformly distributed, the linear system becomes more *numerically stable*.

What *preconditioner* P should we use?

- Standard preconditioners like Jacobi.
- More effective preconditioners. Active area of research.

Is this effective for Poisson solvers?

$$P_{\text{Jacobi}} = \text{diag}(A)$$



Reminder from Variational Case

Poisson equation (with Dirichlet boundary conditions)

$$\Delta f = \operatorname{div} \mathbf{v} \quad \text{over } \Omega, \quad \text{with } f|_{\partial\Omega} = f^*|_{\partial\Omega}$$

After discretization, equivalent to:

$$\begin{bmatrix} D & I & 0 & 0 & 0 & \cdots & 0 \\ I & D & I & 0 & 0 & \cdots & 0 \\ 0 & I & D & I & 0 & \cdots & 0 \\ \vdots & \ddots & \ddots & \ddots & \ddots & \ddots & \vdots \\ 0 & \cdots & 0 & I & D & I & 0 \\ 0 & \cdots & \cdots & 0 & I & D & I \\ 0 & \cdots & \cdots & \cdots & 0 & I & D \end{bmatrix} \cdot \begin{bmatrix} f_1 \\ f_{q_1} \\ \vdots \\ f_{q_2} \\ f_p \\ f_{q_3} \\ \vdots \\ f_{q_4} \\ \vdots \\ f_P \end{bmatrix} = \begin{bmatrix} (\nabla \cdot \mathbf{v})_1 \\ \vdots \\ (\nabla \cdot \mathbf{v})_{q_1} \\ \vdots \\ (\nabla \cdot \mathbf{v})_{q_2} \\ (\nabla \cdot \mathbf{v})_p \\ (\nabla \cdot \mathbf{v})_{q_3} \\ \vdots \\ (\nabla \cdot \mathbf{v})_{q_4} \\ \vdots \\ (\nabla \cdot \mathbf{v})_P \end{bmatrix}$$

Linear system of equations:



$$Af = b$$

Matrix is $P \times P \rightarrow$ billions of entries

WARNING: requires special treatment at the borders
(target boundary values are same as source)



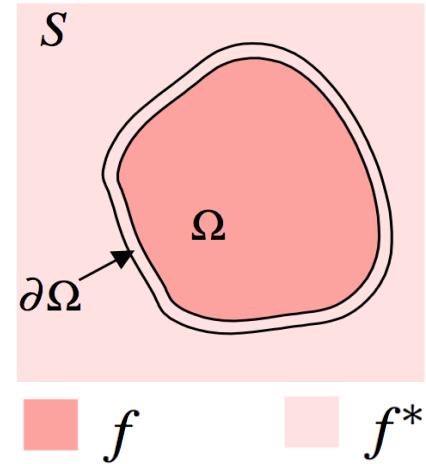
Note: (Dirichlet) Boundary Conditions

- Form a **mask B** that is 0 for pixels that should ***not* be updated** (pixels on $S - \Omega$ and $\partial\Omega$) and 1 otherwise.
- Use **convolution** to perform **Laplacian filtering** over the *entire image*.
- Use (conjugate) gradient descent rules to only update pixels for which the mask is 1. Equivalently, change the update rules to:

$$f^{i+1} = f^i + B\eta^i r^i \quad (\text{gradient descent})$$

$$f^{i+1} = f^i + B\eta^i d^i \quad (\text{conjugate gradient descent})$$

In practice, masking is also required at other steps of (conjugate) gradient descent, to deal with invalid boundaries (e.g., from convolutions).



Poisson Image Editing Examples



Photoshop's “Healing Brush”

Poisson
cloning



Covariant
cloning



Slightly more advanced version of
what we covered here:

- Uses higher-order derivatives



Contrast Problem



Loss of **contrast** when pasting from dark to bright:

- Contrast is a multiplicative property.
- With Poisson blending we are matching linear differences.

Solution: Do blending in **log-domain**.

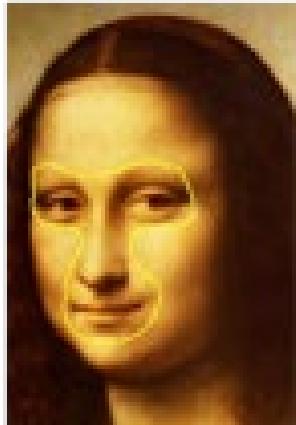


More Blending



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen



originals



copy-paste



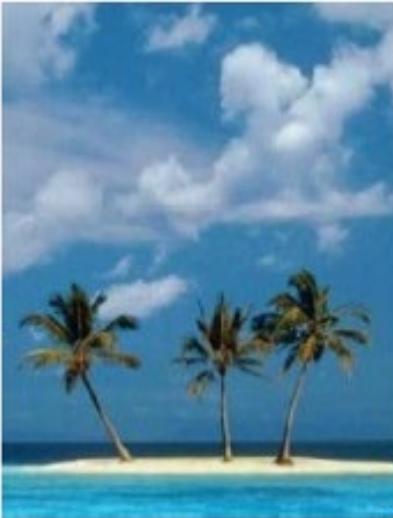
Poisson blending



Blending Transparent Objects



source



destination





香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

Blending Objects with Holes



(a) color-based cutout and paste



(b) seamless cloning



(c) seamless cloning and destination averaged



(d) mixed seamless cloning



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

Editing





Concealment



How would you do this with
Poisson blending?

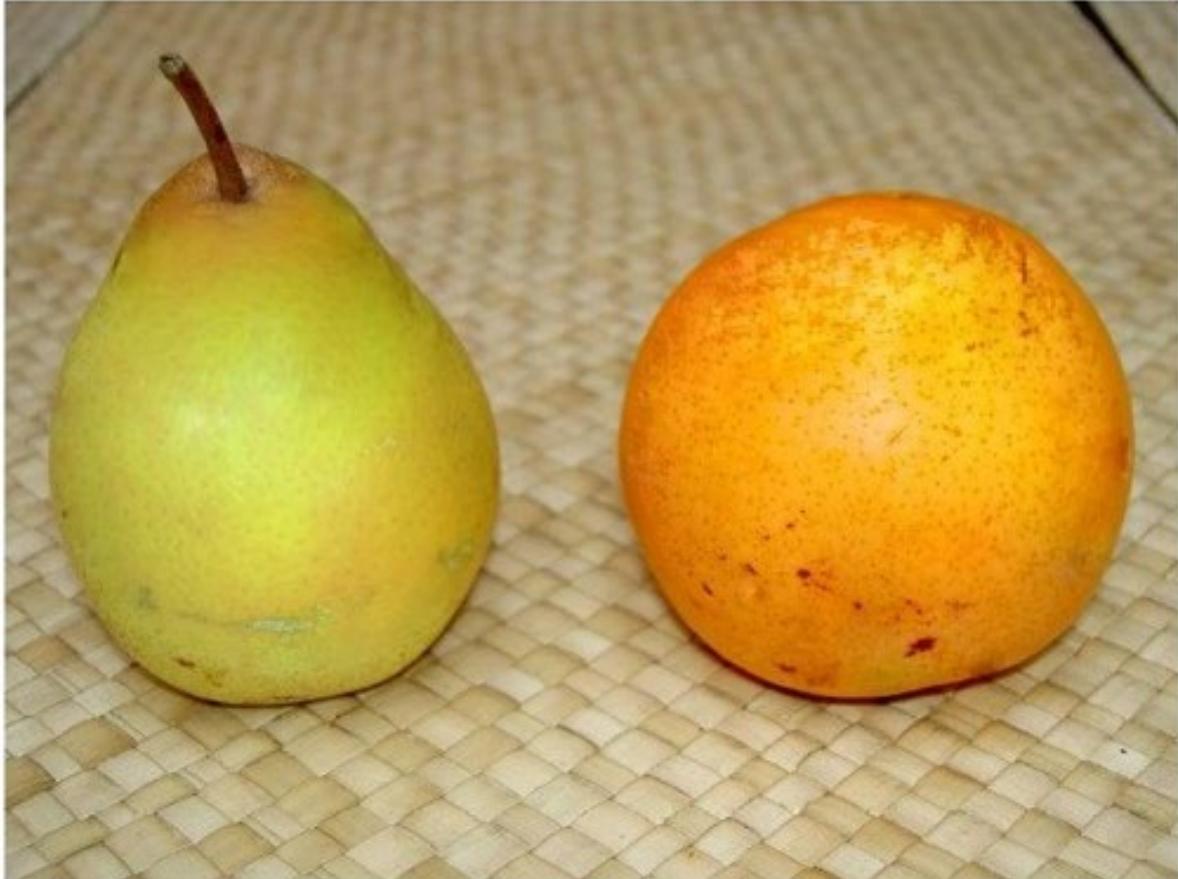
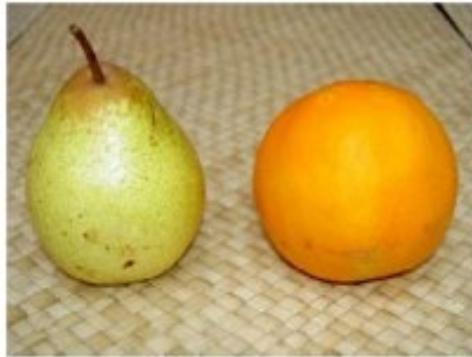
- Insert a copy of the background.



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

Texture Swapping





Special Case: Membrane Interpolation

How would you do this?



Poisson problem

$$\min_f \iint_{\Omega} |\nabla f - \mathbf{v}|^2 \quad \text{with} \quad f|_{\partial\Omega} = f^*|_{\partial\Omega}$$

Laplacian problem

$$\min_f \iint_{\Omega} |\nabla f|^2 \quad \text{with} \quad f|_{\partial\Omega} = f^*|_{\partial\Omega}$$



Entire Suite of Image Editing Tools

GradientShop: A Gradient-Domain Optimization Framework for Image and Video Filtering

Pravin Bhat¹ C. Lawrence Zitnick²

¹University of Washington

Michael Cohen^{1,2} Brian Curless¹

²Microsoft Research



(a) Input image



(b) Saliency-sharpening filter



(c) Pseudo-relighting filter



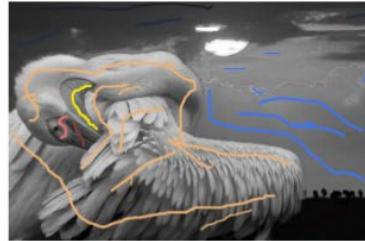
(d) Non-photorealistic rendering filter



(e) Compressed input-image



(f) De-blocking filter

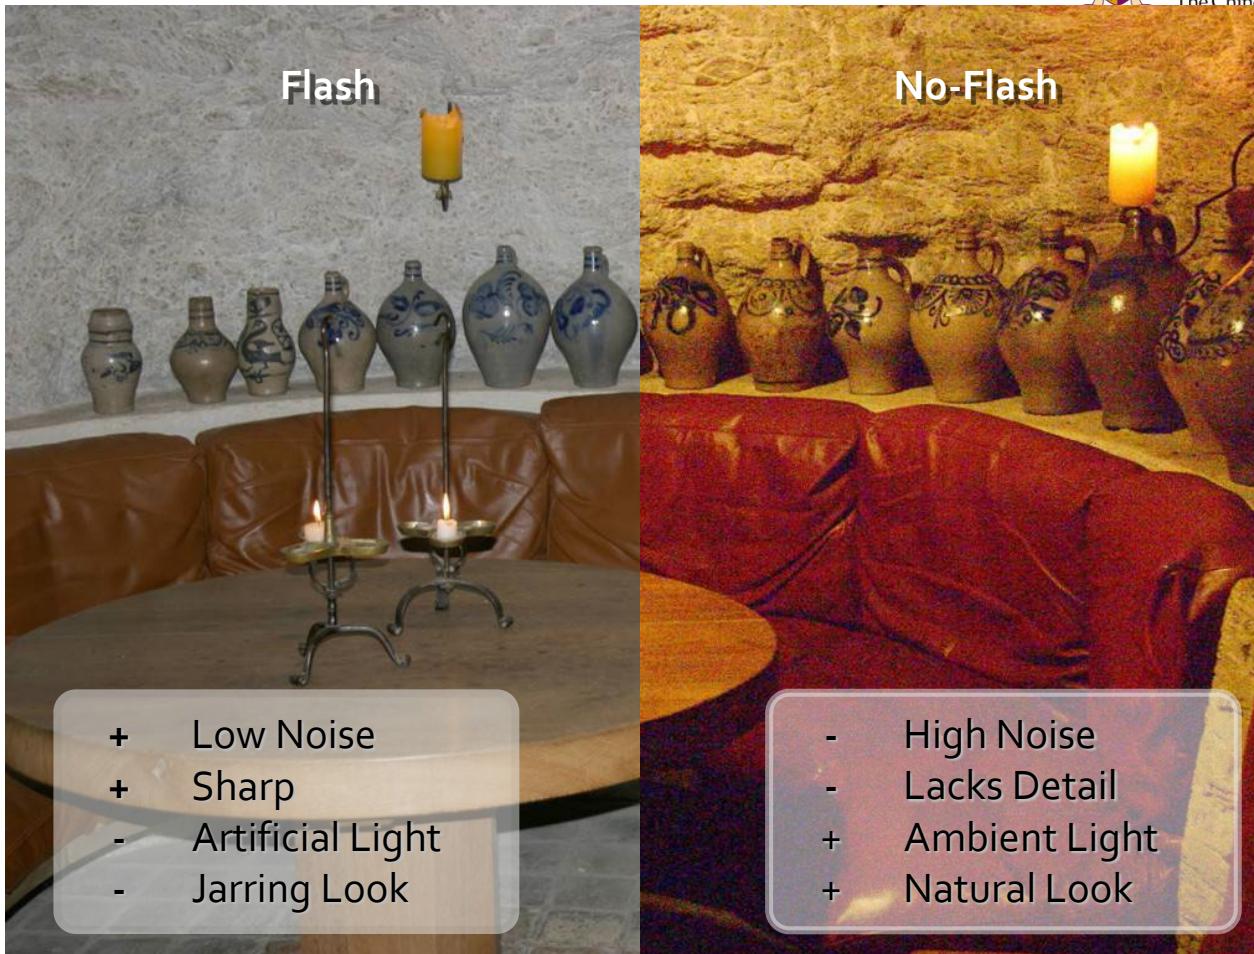


(g) User input for colorization



(h) Colorization filter

Flash/no-flash Photography





香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen



Denoising Result



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen



No-Flash



香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen



Denoising Result



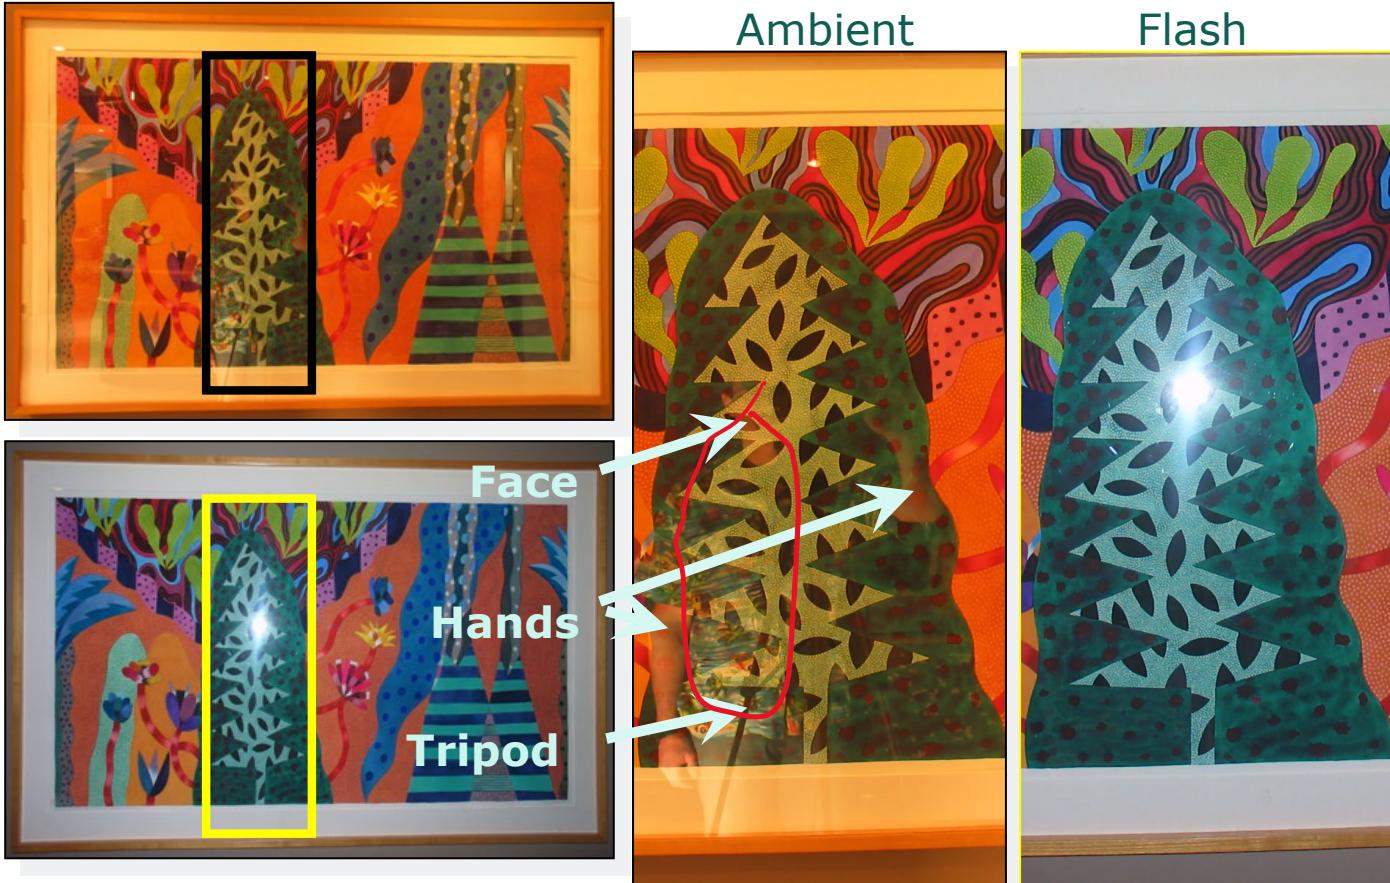
Key Idea

Denoise the no-flash image while maintaining the edge structure of the flash image.

Can we do similar flash/no-flash fusion tasks with gradient-domain processing?

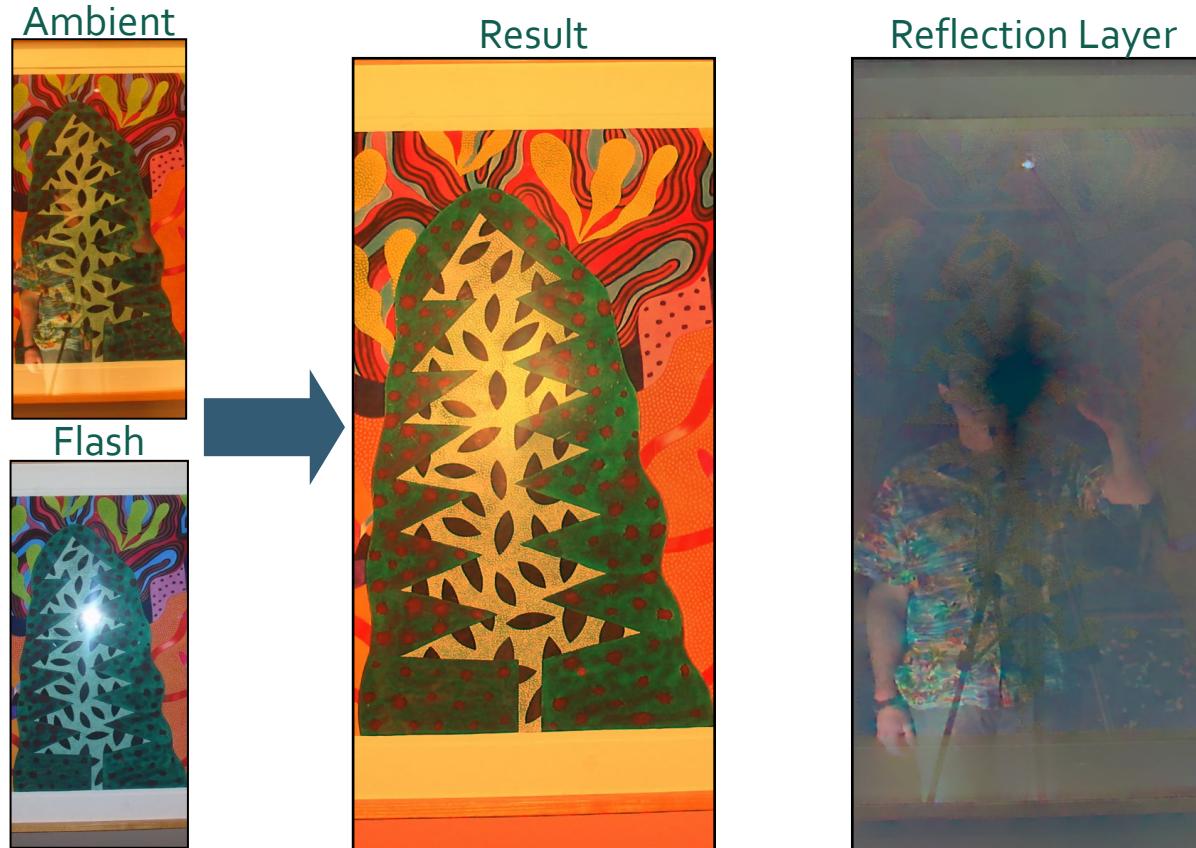


Removing Self-reflections and Hot-spots





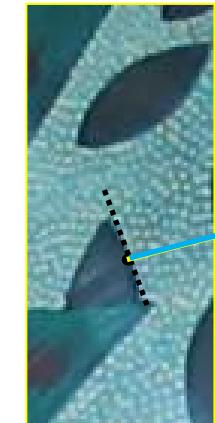
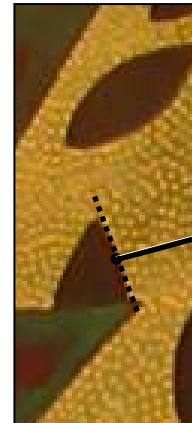
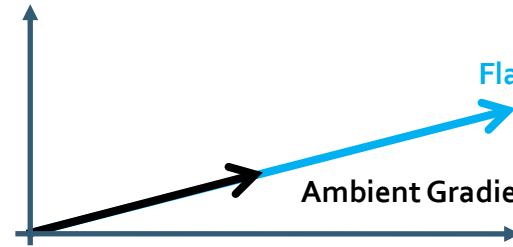
Removing Self-reflections and Hot-spots





Idea: Look at How Gradients are Affected

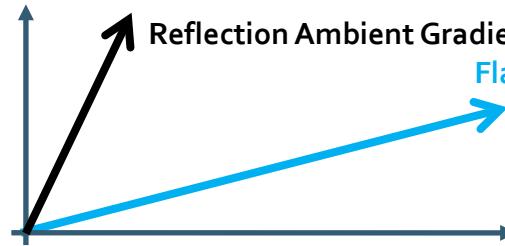
Same gradient vector direction



No reflections

Idea: Look at How Gradients are Affected

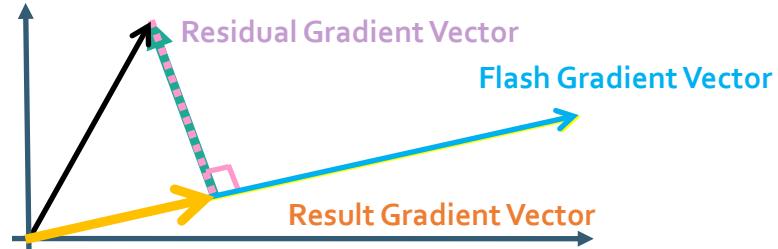
Different gradient vector direction



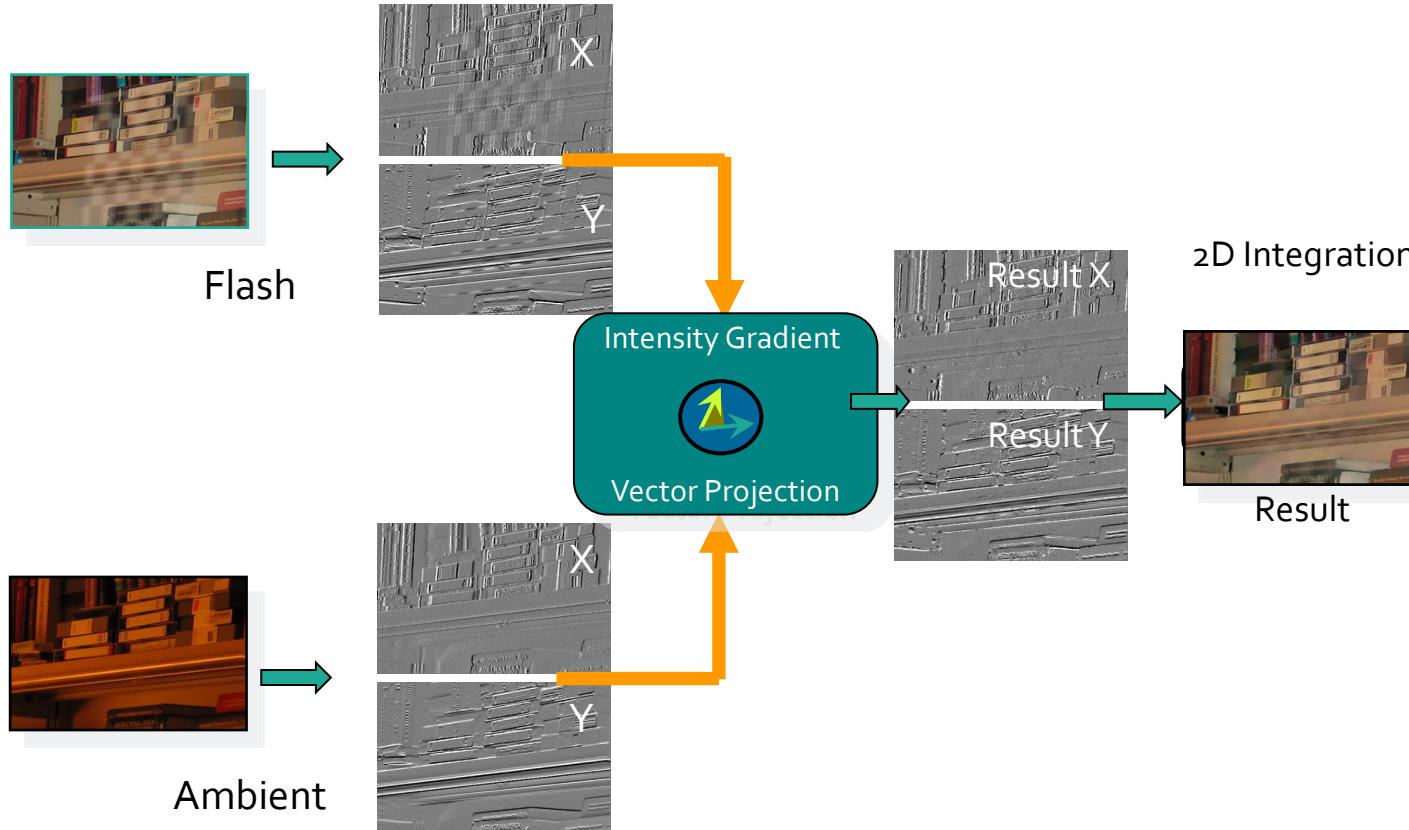
With reflections



Gradient Projections



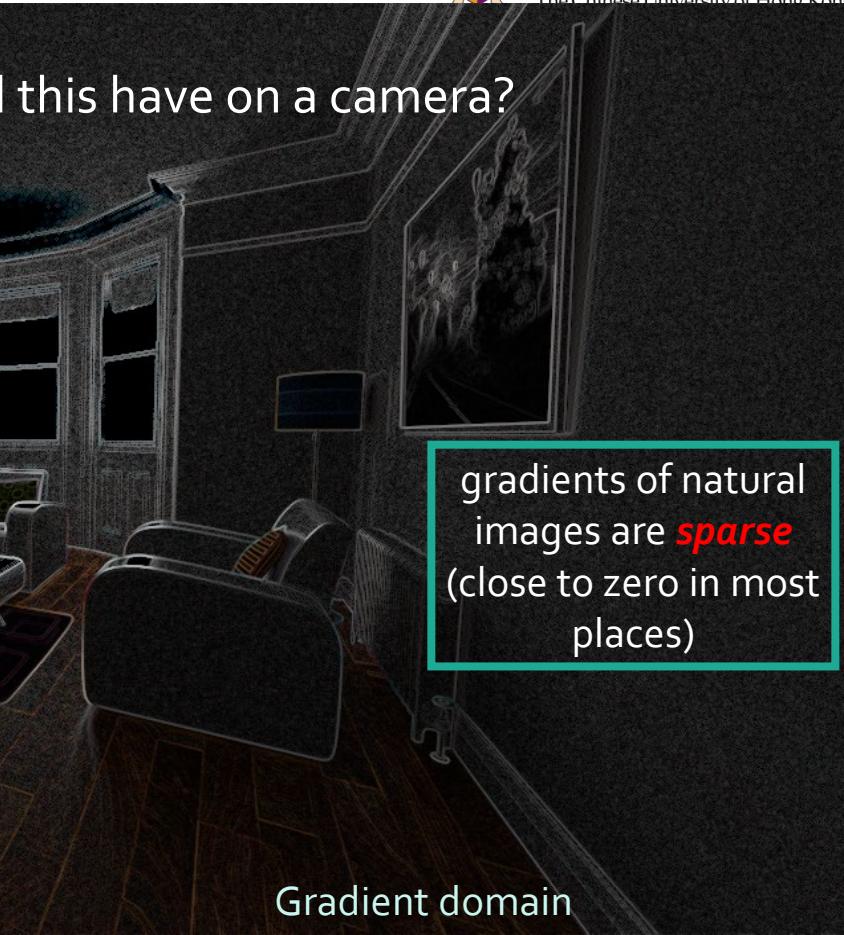
Flash/no-flash with Gradient-domain Processing



Gradient Cameras



What implication would this have on a camera?





Why Gradient Cameras?

Why I want a Gradient Camera

Jack Tumblin
Northwestern University
jet@cs.northwestern.edu

Amit Agrawal
University of Maryland
aagrawal@umd.edu

Ramesh Raskar
MERL
raskar@merl.com

Why a gradient camera?

- Faster frame rate, only very few pixels (gradient).
- Higher dynamic range, if also combined with logarithmic gradients.

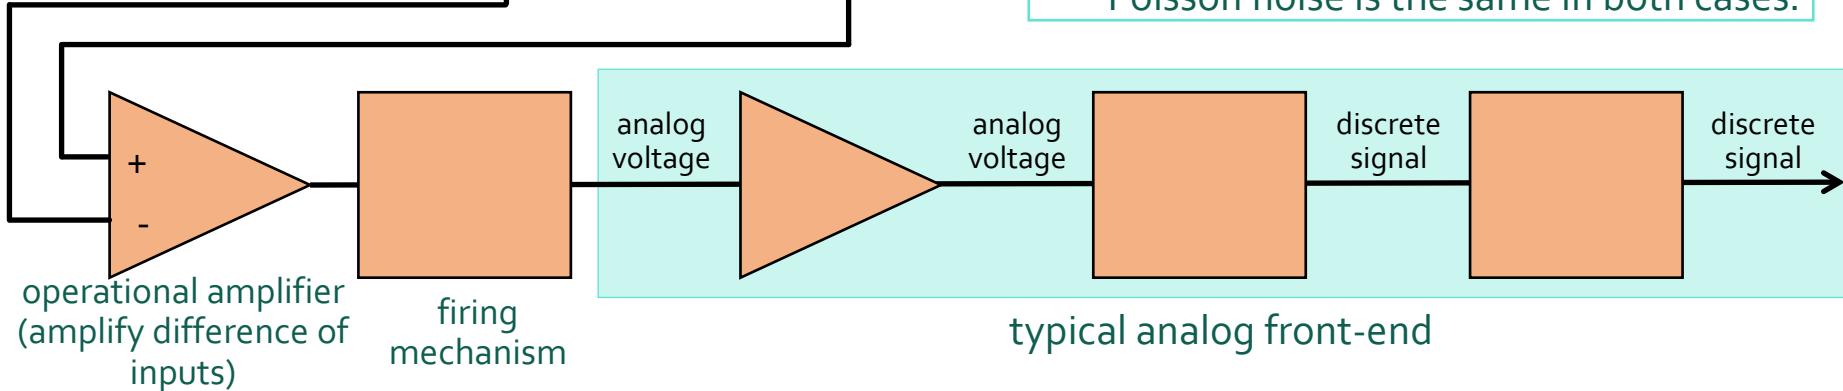
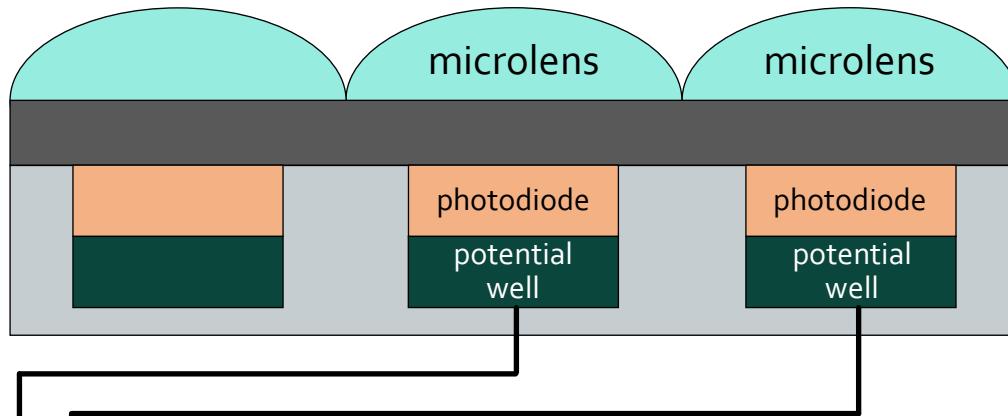
Can you directly display the measurements of such a camera?

- Use a Poisson solver to reconstruct the image from the measured gradients.

How to build a gradient camera?



Change the Sensor



Any disadvantages?

- Spatial resolution is reduced by 2x.
- Photosensitive area is reduced.

Why is this better than computing gradients in post-processing?

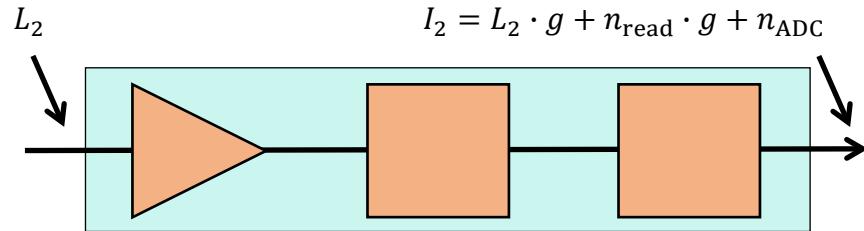
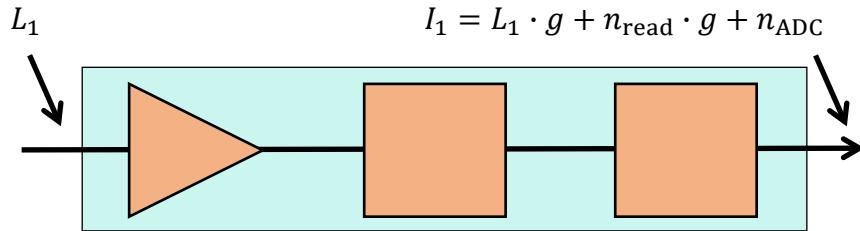
- Additive noise is reduced.
- Acquisition is faster thanks to the firing mechanism and sparsity of edges.

What about Poisson noise?

- Poisson noise is the same in both cases.



Noise of Diff Sensor



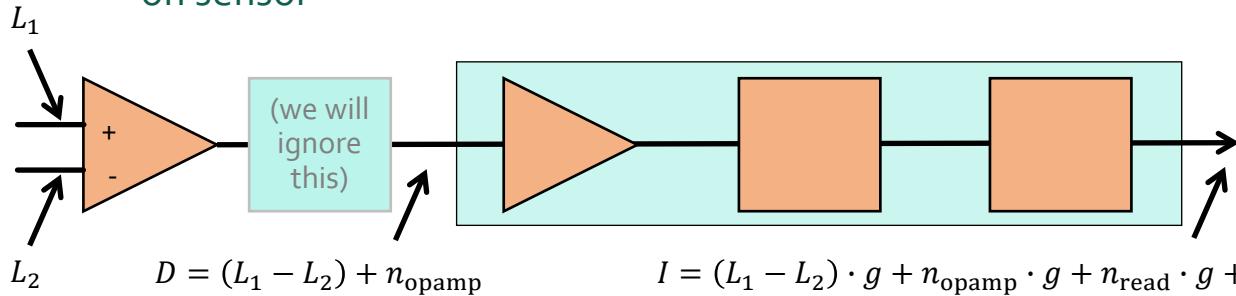
Digital subtraction in post-processing

$$\sigma(I_1 - I_2)^2 = \sigma(L_1 - L_2)^2 + 2 \cdot \sigma_{\text{read}}^2 \cdot g^2 + 2 \cdot \sigma_{\text{ADC}}^2$$

which variance is better?

Analog subtraction on sensor

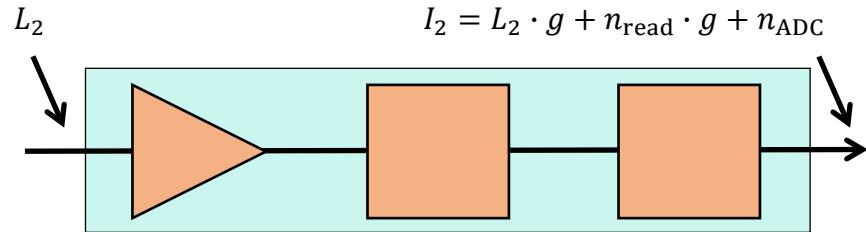
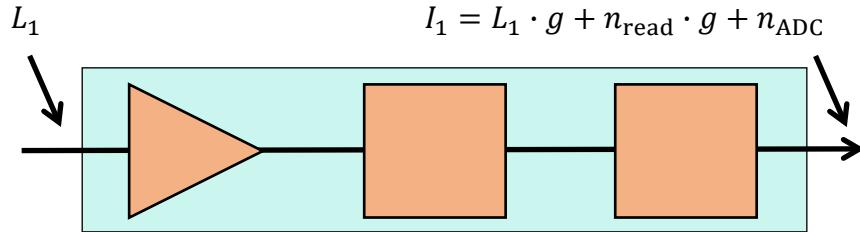
$$\sigma(I)^2 = \sigma(L_1 - L_2)^2 + \sigma_{\text{opamp}}^2 \cdot g^2 + \sigma_{\text{read}}^2 \cdot g^2 + \sigma_{\text{ADC}}^2$$



$$\begin{aligned}L_1 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_1 + D)) \\L_2 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_2 + D)) \\n_{\text{opamp}} &\sim \text{Normal}(0, \sigma_{\text{opamp}}) \\n_{\text{read}} &\sim \text{Normal}(0, \sigma_{\text{read}}) \\n_{\text{ADC}} &\sim \text{Normal}(0, \sigma_{\text{ADC}})\end{aligned}$$



Noise of Diff Sensor



Digital subtraction in post-processing

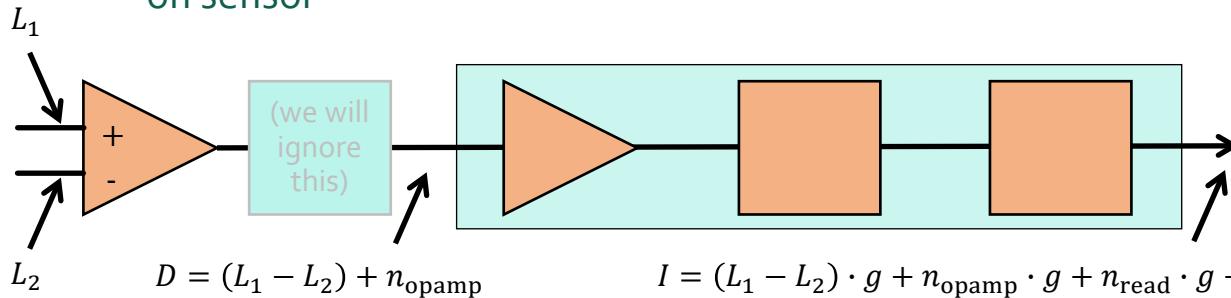
$$\sigma(I_1 - I_2)^2 = \boxed{\sigma(L_1 - L_2)^2} + \boxed{2 \cdot \sigma_{\text{read}}^2 \cdot g^2 + 2 \cdot \sigma_{\text{ADC}}^2}$$

terms related to Poisson noise are the same

additive noise is reduced if opamp is well-designed

Analog subtraction on sensor

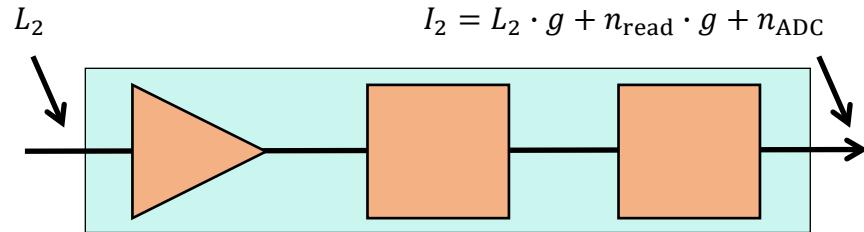
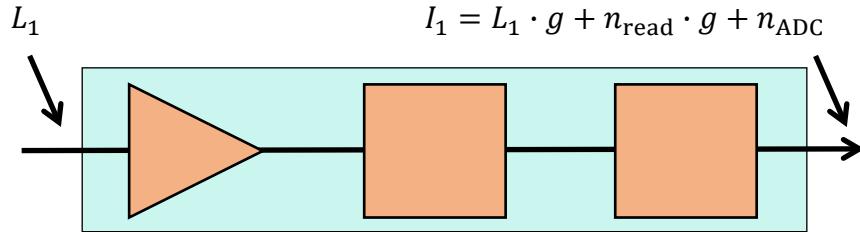
$$\sigma(I)^2 = \boxed{\sigma(L_1 - L_2)^2} + \boxed{\sigma_{\text{opamp}}^2 \cdot g^2 + \sigma_{\text{read}}^2 \cdot g^2 + \sigma_{\text{ADC}}^2}$$



$$\begin{aligned} L_1 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_1 + D)) \\ L_2 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_2 + D)) \\ n_{\text{opamp}} &\sim \text{Normal}(0, \sigma_{\text{opamp}}) \\ n_{\text{read}} &\sim \text{Normal}(0, \sigma_{\text{read}}) \\ n_{\text{ADC}} &\sim \text{Normal}(0, \sigma_{\text{ADC}}) \end{aligned}$$



Noise of Diff Sensor

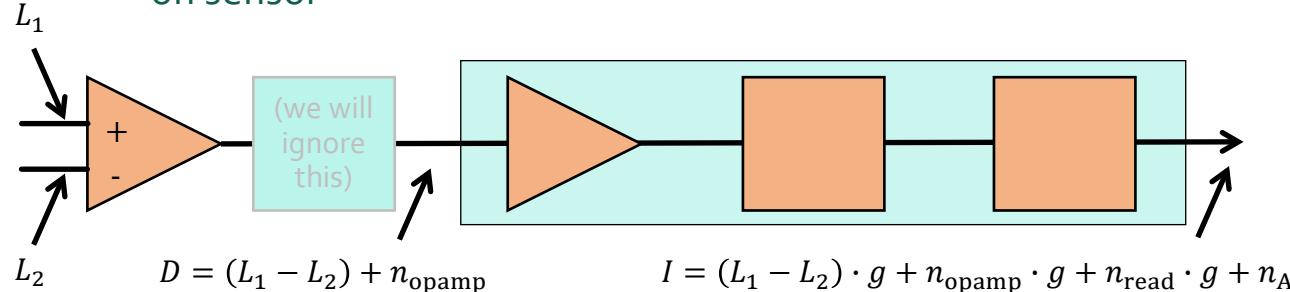


Digital subtraction in post-processing

$$\sigma(I_1 - I_2)^2 = \boxed{\sigma(L_1 - L_2)^2} + 2 \cdot \sigma_{\text{read}}^2 \cdot g^2 + 2 \cdot \sigma_{\text{ADC}}^2$$

what is the distribution of the difference $L_1 - L_2$?

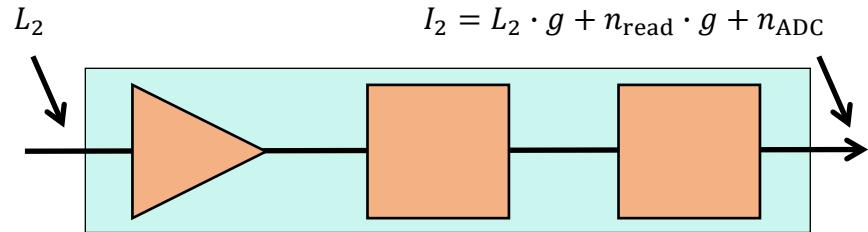
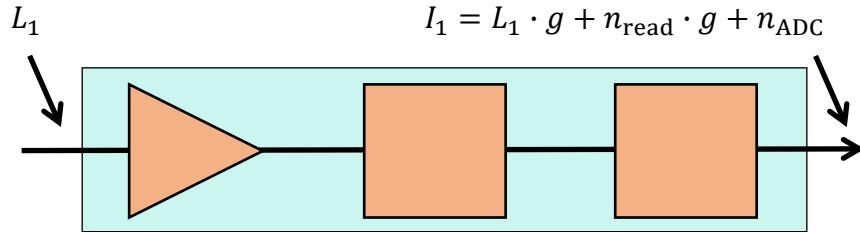
Analog subtraction on sensor



$$\begin{aligned} L_1 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_1 + D)) \\ L_2 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_2 + D)) \\ n_{\text{opamp}} &\sim \text{Normal}(0, \sigma_{\text{opamp}}) \\ n_{\text{read}} &\sim \text{Normal}(0, \sigma_{\text{read}}) \\ n_{\text{ADC}} &\sim \text{Normal}(0, \sigma_{\text{ADC}}) \end{aligned}$$



Noise of Diff Sensor

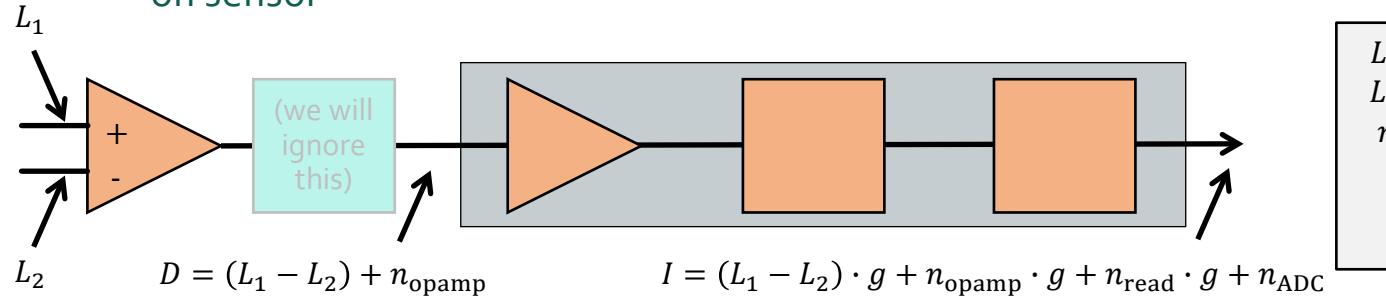


Digital subtraction in post-processing

$$\sigma(I_1 - I_2)^2 = \boxed{\sigma(L_1 - L_2)^2} + 2 \cdot \sigma_{\text{read}}^2 \cdot g^2 + 2 \cdot \sigma_{\text{ADC}}^2$$

$$L_1 - L_2 \sim \text{Skellam}\left(t \cdot a \cdot (\Phi_1 - \Phi_2), t \cdot (a \cdot (\Phi_1 + \Phi_2) + 2 \cdot D)\right)$$

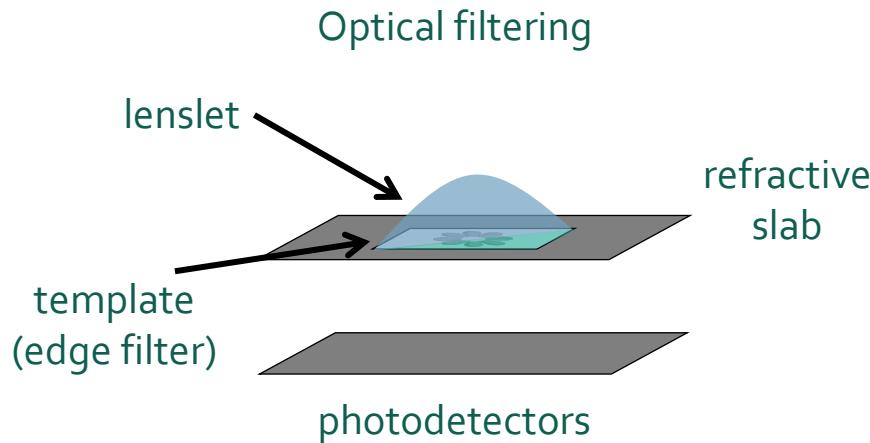
Analog subtraction on sensor



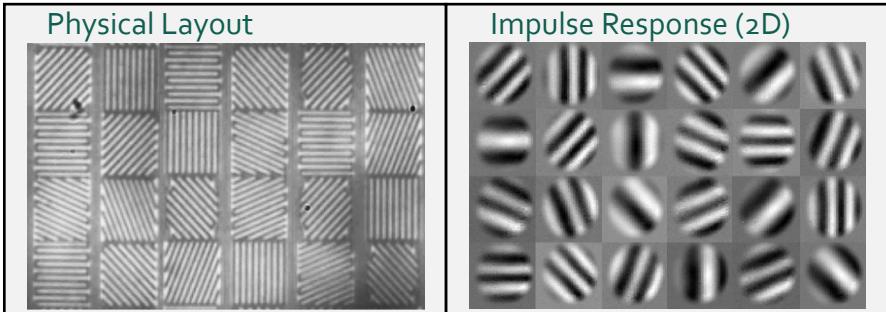
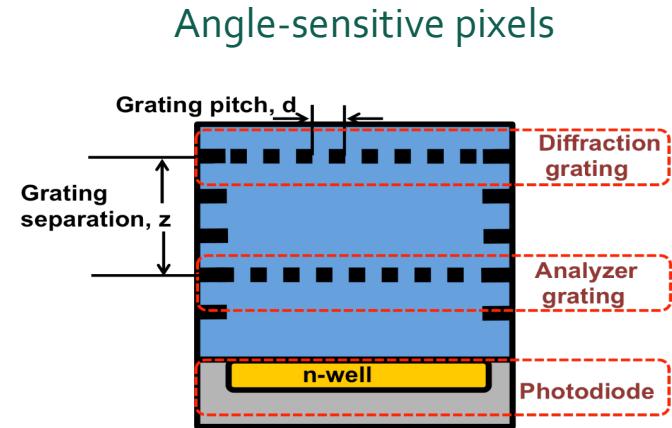
$$\begin{aligned} L_1 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_1 + D)) \\ L_2 &\sim \text{Poisson}(t \cdot (a \cdot \Phi_2 + D)) \\ n_{\text{opamp}} &\sim \text{Normal}(0, \sigma_{\text{opamp}}) \\ n_{\text{read}} &\sim \text{Normal}(0, \sigma_{\text{read}}) \\ n_{\text{ADC}} &\sim \text{Normal}(0, \sigma_{\text{ADC}}) \end{aligned}$$



Change the Optics



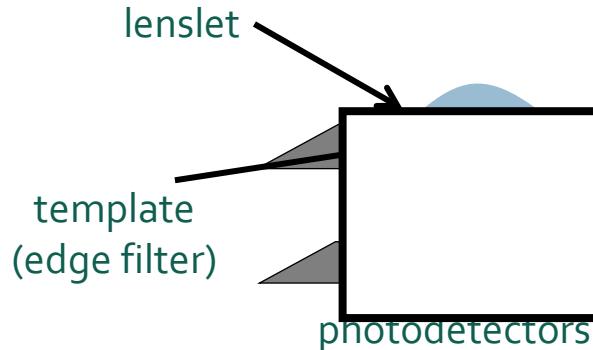
resulting image





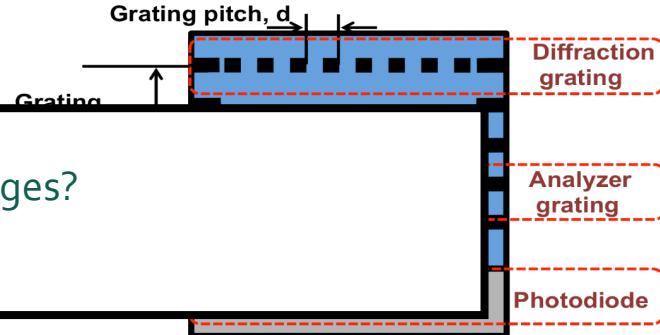
Change the Optics

Optical filtering



refractive

Angle-sensitive pixels

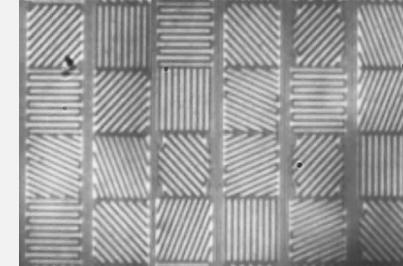


Any disadvantages?

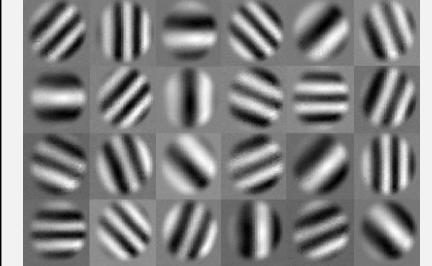


resulting image

Physical Layout



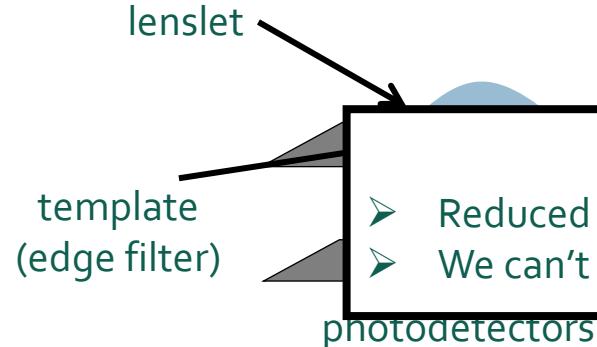
Impulse Response (2D)





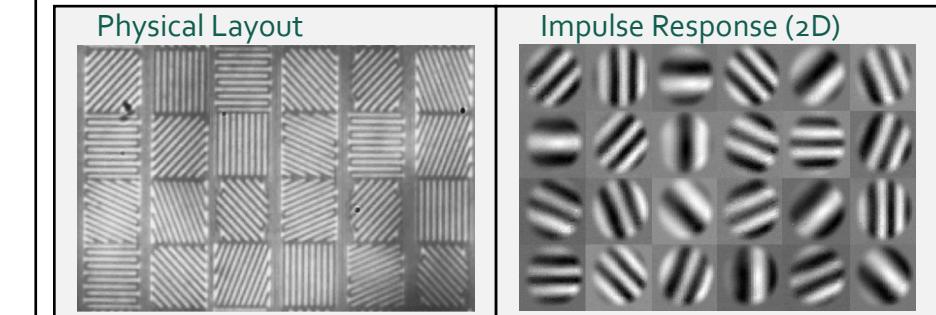
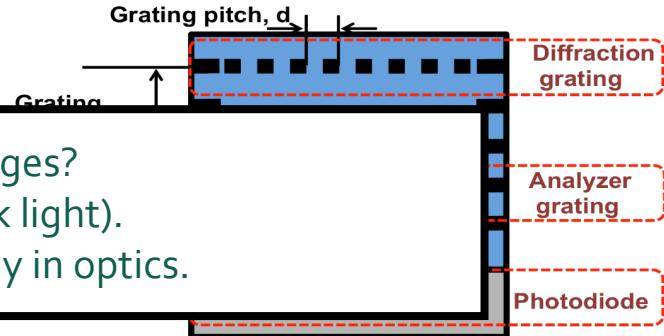
Change the Optics

Optical filtering



resulting image

Angle-sensitive pixels





Why Gradient Cameras?

Why I want a Gradient Camera

Jack Tumblin
Northwestern University
jet@cs.northwestern.edu

Amit Agrawal
University of Maryland
aagrawal@umd.edu

Ramesh Raskar
MERL
raskar@merl.com

Why a gradient camera?

- Faster frame rate, only very few pixels (gradient).
- Higher dynamic range, if also combined with logarithmic gradients.

Can you directly display the measurements of such a camera?

- Use a Poisson solver to reconstruct the image from the measured gradients.

How to build a gradient camera?

- Change the sensor.
- Change the optics.

Temporal Gradients



Event-based camera:

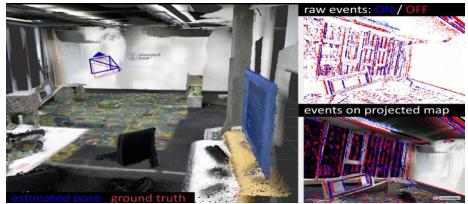
<https://www.youtube.com/watch?v=kPCZESVfHoQ>

High-speed output on a quadcopter:

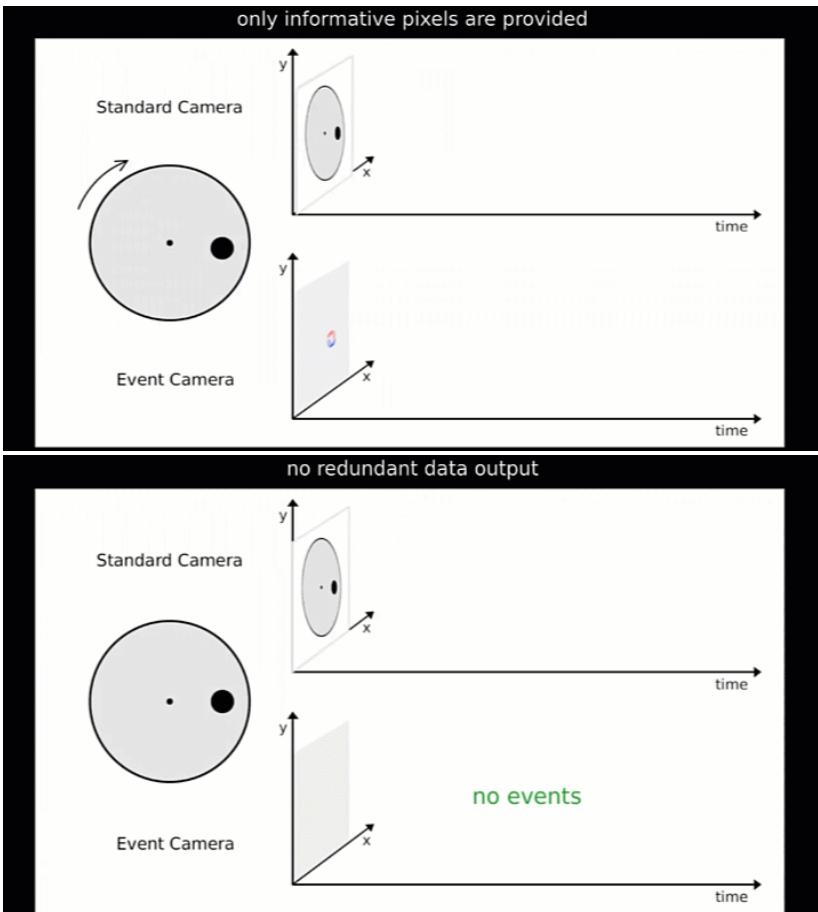
<https://www.youtube.com/watch?v=LauQ6LWTkxM>

Simulator:

<http://rpg.ifi.uzh.ch/esim>



event-based cameras (a.k.a. dynamic vision sensors, or DVS)



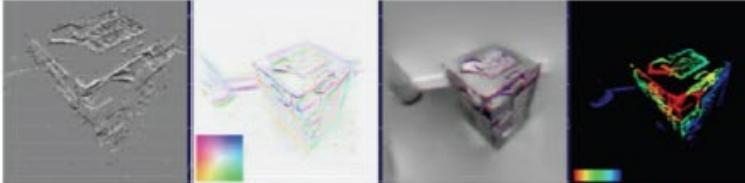


香港中文大學(深圳)

The Chinese University of Hong Kong, Shenzhen

Event Camera in Robotics and Cision

box
(indoor)



kitchen
(indoor)



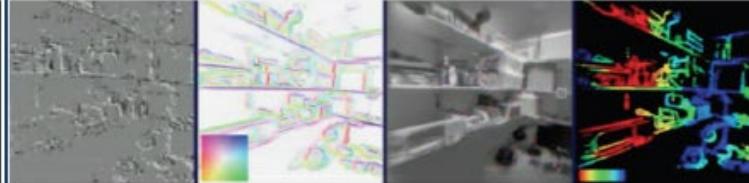
objects
(indoor)



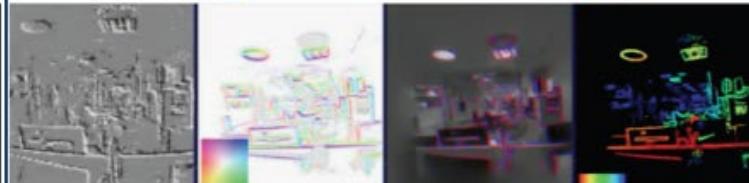
bikes
(outdoor)



shelves
(indoor)



office
(indoor)



hallway
(outdoor)



statue
(outdoor)



event
visualisation

gradient
estimation

intensity
reconstruction

depth
estimation

event
visualisation

gradient
estimation

intensity
reconstruction

depth
estimation



References

Basic reading:

- Szeliski textbook, Sections 3.13, 3.5.5, 9.3.4, 10.4.3.
- Pérez et al., "Poisson Image Editing," SIGGRAPH 2003.
 The original Poisson image editing paper.
- Agrawal and Raskar, "Gradient Domain Manipulation Techniques in Vision and Graphics," ICCV 2007 course, <http://www.amitkagrwal.com/ICCV2007Course/>
 A great resource (entire course!) for gradient-domain image processing.
- Agrawal et al., "Removing Photography Artifacts Using Gradient Projection and Flash-Exposure Sampling," SIGGRAPH 2005.
 A paper on photography with flash and no-flash pairs, using gradient-domain image processing.

Additional reading:

- Georgiev, "Covariant Derivatives and Vision," ECCV 2006.
 An paper from Adobe on the version of Poisson blending implemented in Photoshop's "healing brush".
- Elder and Goldberg, "Image editing in the contour domain", PAMI 2001.
 One of the very first papers discussing gradient-domain image processing.
- Frankot and Chellappa, "A method for enforcing integrability in shape from shading algorithms," PAMI 1988.
- Bhat et al., "Fourier Analysis of the 2D Screened Poisson Equation for Gradient Domain Problems," ECCV 2008.
 A couple of papers discussing the (Fourier) basis projection approach for solving the Poisson integration problem.
- Agrawal et al., "What Is the Range of Surface Reconstructions from a Gradient Field?", ECCV 2006.
- Qu  au et al., "Normal Integration: A Survey," JMIV 2017.
 Two papers reviewing various gradient (and surface normal) integration techniques, including Poisson solvers.
- Szeliski, "Locally adapted hierarchical basis preconditioning," SIGGRAPH 2006.
- Krishnan and Szeliski, "Multigrid and multilevel preconditioners for computational photography," SIGGRAPH 2011.
- Krishnan et al., "Efficient Preconditioning of Laplacian Matrices for Computer Graphics," SIGGRAPH 2013.
 A few well-known references on multi-grid and preconditioning techniques for accelerating the Poisson solver, with a specific focus on computational photography applications..
- Shewchuk, "An Introduction to the Conjugate Gradient Method Without the Agonizing Pain," CMU TR 1994, <http://www.cs.cmu.edu/~quake-papers/painless-conjugate-gradient.pdf>
 A great reference on (preconditioned) conjugate gradient solvers for large linear systems.
- Briggs et al., "A multigrid tutorial," SIAM 2000.
 A great reference book on multi-grid approaches.
- Bhat et al., "GradientShop: A Gradient-Domain Optimization Framework for Image and Video Filtering," TOG 2010.
 A paper describing gradient-domain processing as a general image processing paradigm, which can be used for a broad set of applications beyond blending.
- Krishnan and Fergus, "Dark Flash Photography," SIGGRAPH 2009.
 A paper proposing doing flash/no-flash photography using infrared flash lights.
- Kazhdan et al., "Poisson surface reconstruction," SGP 2006.
- Kazhdan and Hoppe, "Screened Poisson surface reconstruction," TOG 2013.
 Two papers discussing Poisson problems for reconstructing meshes from point clouds and normals. This is arguably the most commonly used surface reconstruction algorithm.
- Lehtinen et al., "Gradient-domain metropolis light transport," SIGGRAPH 2013.
- Kettunen et al., "Gradient-domain path tracing," SIGGRAPH 2015.
- Hua et al., "Light transport simulation in the gradient domain," SIGGRAPH Asia 2018 course, http://beltegeuse.s3-website-ap-northeast-1.amazonaws.com/research/2018_GradientCourse/
 In addition to *editing* images in the gradient-domain, we can *render* them directly in the gradient-domain.
- Tumblin et al., "Why I want a gradient camera?" CVPR 2005.
 We can even directly *measure* images in the gradient domain, using so-called *gradient cameras*.
- Callenberg et al., "Snapshot difference imaging using correlation time-of-flight sensors," SIGGRAPH Asia 2017.
 A form of camera with differential pixels.
- Koppal et al., "Toward wide-angle microvision sensors", PAMI 2013.
 Gradient cameras using optical filtering.
- Chen et al., "ASP vision: Optically computing the first layer of convolutional neural networks using angle sensitive pixels," CVPR 2016.
 Gradient cameras using angle-sensitive pixels.
- Kim et al., "Real-time 3D reconstruction and 6-DoF tracking with an event camera," ECCV 2016.
 A paper on using event-based cameras for computer vision applications in very fast frame rates (best paper award at ECCV 2016!).



Today's Topic

- Efficient Poisson Solver
- Poisson Image Editing
- Flash/no-flash Photography
- Gradient Cameras
 - Diff Sensor
 - Diff Optics
 - Diff Time



Thank You!



Qilin Sun (孙启霖)

School of Data Science

The Chinese University of Hong Kong, Shenzhen